

experiments

April 7, 2025

1 Feature Transformation Using Sklearn with ANN

```
[1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
import pickle
```

```
[2]: # Load the dataset
data = pd.read_csv('../Datasets/Churn_Modelling.csv')
data.head()
```

```
[2]:  RowNumber  CustomerId  Surname  CreditScore  Geography  Gender  Age  \
0         1    15634602   Hargrave         619     France  Female   42
1         2    15647311     Hill         608     Spain  Female   41
2         3    15619304     Onio         502     France  Female   42
3         4    15701354     Boni         699     France  Female   39
4         5    15737888  Mitchell         850     Spain  Female   43
```

```
      Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  \
0         2     0.00             1           1              1
1         1  83807.86             1           0              1
2         8 159660.80             3           1              0
3         1     0.00             2           0              0
4         2 125510.82             1           1              1
```

```
      EstimatedSalary  Exited
0         101348.88       1
1         112542.58       0
2         113931.57       1
3          93826.63       0
4          79084.10       0
```

```
[3]: # Preprocess the data
## Drop irrelevant columns

data = data.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1)
data
```

```
[3]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	\
0	619	France	Female	42	2	0.00	1	
1	608	Spain	Female	41	1	83807.86	1	
2	502	France	Female	42	8	159660.80	3	
3	699	France	Female	39	1	0.00	2	
4	850	Spain	Female	43	2	125510.82	1	
...	
9995	771	France	Male	39	5	0.00	2	
9996	516	France	Male	35	10	57369.61	1	
9997	709	France	Female	36	7	0.00	1	
9998	772	Germany	Male	42	3	75075.31	2	
9999	792	France	Female	28	4	130142.79	1	

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1
1	0	1	112542.58	0
2	1	0	113931.57	1
3	0	0	93826.63	0
4	1	1	79084.10	0
...
9995	1	0	96270.64	0
9996	1	1	101699.77	0
9997	0	1	42085.58	1
9998	1	0	92888.52	1
9999	1	0	38190.78	0

[10000 rows x 11 columns]

```
[4]: ## Encode categorical variables
label_encoder_gender = LabelEncoder() # we use this for converting the char
↳data to int data
data['Gender']= label_encoder_gender.fit_transform(data['Gender'])
data
```

```
[4]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	\
0	619	France	0	42	2	0.00	1	
1	608	Spain	0	41	1	83807.86	1	
2	502	France	0	42	8	159660.80	3	
3	699	France	0	39	1	0.00	2	
4	850	Spain	0	43	2	125510.82	1	
...	
9995	771	France	1	39	5	0.00	2	
9996	516	France	1	35	10	57369.61	1	
9997	709	France	0	36	7	0.00	1	
9998	772	Germany	1	42	3	75075.31	2	
9999	792	France	0	28	4	130142.79	1	

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1
1	0	1	112542.58	0
2	1	0	113931.57	1
3	0	0	93826.63	0
4	1	1	79084.10	0
...
9995	1	0	96270.64	0
9996	1	1	101699.77	0
9997	0	1	42085.58	1
9998	1	0	92888.52	1
9999	1	0	38190.78	0

[10000 rows x 11 columns]

```
[5]: ## Onehot encode "Geography"
from sklearn.preprocessing import OneHotEncoder
onehot_encoder_geo=OneHotEncoder()
geo_encoder= onehot_encoder_geo.fit_transform(data[['Geography']])
geo_encoder
```

```
[5]: <Compressed Sparse Row sparse matrix of dtype 'float64'
      with 10000 stored elements and shape (10000, 3)>
```

```
[6]: geo_encoder.toarray()
```

```
[6]: array([[1., 0., 0.],
          [0., 0., 1.],
          [1., 0., 0.],
          ...,
          [1., 0., 0.],
          [0., 1., 0.],
          [1., 0., 0.]])
```

```
[7]: onehot_encoder_geo.get_feature_names_out(['Geography'])
```

```
[7]: array(['Geography_France', 'Geography_Germany', 'Geography_Spain'],
      dtype=object)
```

```
[8]: geo_encoder_df=pd.DataFrame(geo_encoder.toarray(), columns=onehot_encoder_geo.
    ↪get_feature_names_out(['Geography']))
geo_encoder_df
```

	Geography_France	Geography_Germany	Geography_Spain
0	1.0	0.0	0.0
1	0.0	0.0	1.0
2	1.0	0.0	0.0
3	1.0	0.0	0.0

4	0.0	0.0	1.0
...
9995	1.0	0.0	0.0
9996	1.0	0.0	0.0
9997	1.0	0.0	0.0
9998	0.0	1.0	0.0
9999	1.0	0.0	0.0

[10000 rows x 3 columns]

```
[9]: ## Combine one hot encoder columns with the original data
data = pd.concat([data.drop('Geography', axis=1), geo_encoder_df], axis=1)
data.head()
```

```
[9]:   CreditScore  Gender  Age  Tenure   Balance  NumOfProducts  HasCrCard  \
0         619      0   42      2     0.00           1           1
1         608      0   41      1  83807.86           1           0
2         502      0   42      8 159660.80           3           1
3         699      0   39      1     0.00           2           0
4         850      0   43      2 125510.82           1           1
```

	IsActiveMember	EstimatedSalary	Exited	Geography_France	\
0	1	101348.88	1	1.0	
1	1	112542.58	0	0.0	
2	0	113931.57	1	1.0	
3	0	93826.63	0	1.0	
4	1	79084.10	0	0.0	

	Geography_Germany	Geography_Spain
0	0.0	0.0
1	0.0	1.0
2	0.0	0.0
3	0.0	0.0
4	0.0	1.0

```
[10]: ## Save the encoders and standeredScaler
with open('label_encoder_gender.pkl', 'wb') as file:
    pickle.dump(label_encoder_gender, file)

with open('onehot_encoder_geo.pkl', 'wb') as file:
    pickle.dump(onehot_encoder_geo, file)
```

```
[11]: data.head()
```

```
[11]:   CreditScore  Gender  Age  Tenure   Balance  NumOfProducts  HasCrCard  \
0         619      0   42      2     0.00           1           1
1         608      0   41      1  83807.86           1           0
```

2	502	0	42	8	159660.80	3	1
3	699	0	39	1	0.00	2	0
4	850	0	43	2	125510.82	1	1

	IsActiveMember	EstimatedSalary	Exited	Geography_France \
0	1	101348.88	1	1.0
1	1	112542.58	0	0.0
2	0	113931.57	1	1.0
3	0	93826.63	0	1.0
4	1	79084.10	0	0.0

	Geography_Germany	Geography_Spain
0	0.0	0.0
1	0.0	1.0
2	0.0	0.0
3	0.0	0.0
4	0.0	1.0

```
[12]: ## Divide the dataset into independent and dependent features
X = data.drop('Exited', axis=1)
y = data['Exited']

## Split the data in the training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

## Scale these features
scaler = StandardScaler()
X_train= scaler.fit_transform(X_train)
X_test= scaler.transform(X_test)
```

```
[13]: X_train
```

```
[13]: array([[ 0.35649971,  0.91324755, -0.6557859 , ...,  1.00150113,
           -0.57946723, -0.57638802],
           [-0.20389777,  0.91324755,  0.29493847, ..., -0.99850112,
            1.72572313, -0.57638802],
           [-0.96147213,  0.91324755, -1.41636539, ..., -0.99850112,
            -0.57946723,  1.73494238],
           ...,
           [ 0.86500853, -1.09499335, -0.08535128, ...,  1.00150113,
            -0.57946723, -0.57638802],
           [ 0.15932282,  0.91324755,  0.3900109 , ...,  1.00150113,
            -0.57946723, -0.57638802],
           [ 0.47065475,  0.91324755,  1.15059039, ..., -0.99850112,
            1.72572313, -0.57638802]])
```

```
[14]: with open('scaler.pkl', 'wb') as file:
      pickle.dump(scaler, file)
```

```
[15]: data
```

```
[15]:
```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	619	0	42	2	0.00	1	1	
1	608	0	41	1	83807.86	1	0	
2	502	0	42	8	159660.80	3	1	
3	699	0	39	1	0.00	2	0	
4	850	0	43	2	125510.82	1	1	
...	
9995	771	1	39	5	0.00	2	1	
9996	516	1	35	10	57369.61	1	1	
9997	709	0	36	7	0.00	1	0	
9998	772	1	42	3	75075.31	2	1	
9999	792	0	28	4	130142.79	1	1	

	IsActiveMember	EstimatedSalary	Exited	Geography_France	\
0	1	101348.88	1	1.0	
1	1	112542.58	0	0.0	
2	0	113931.57	1	1.0	
3	0	93826.63	0	1.0	
4	1	79084.10	0	0.0	
...	
9995	0	96270.64	0	1.0	
9996	1	101699.77	0	1.0	
9997	1	42085.58	1	1.0	
9998	0	92888.52	1	0.0	
9999	0	38190.78	0	1.0	

	Geography_Germany	Geography_Spain
0	0.0	0.0
1	0.0	1.0
2	0.0	0.0
3	0.0	0.0
4	0.0	1.0
...
9995	0.0	0.0
9996	0.0	0.0
9997	0.0	0.0
9998	1.0	0.0
9999	0.0	0.0

```
[10000 rows x 13 columns]
```

2 Step by Step Training with ANN With Optimizer and Loss function

```
[16]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.callbacks import EarlyStopping, TensorBoard
import datetime
```

2.1 Build ANN Model

```
[17]: (X_train.shape[1],)
```

```
[17]: (12,)
```

So our ANN model is also called as Sequential model so in case of our project we have: - total 12 input data. - first hidden layer of 64 nodes or neurons. - second hidden layer of 32 nodes or neurons. - And the output layer have 1 neuron.

we use the Dense to create the layer which contain the neurons.

```
[18]: model = Sequential([
    Dense(64, activation='relu', input_shape=(X_train.shape[1],)), ## HL-1
    ↪Connected with input layer
    Dense(32, activation='relu'), ## HL-2
    Dense(1, activation='sigmoid') ## Output layer
])
```

```
c:\Users\Uditya\Desktop\Deep Learning\Deep Learning for Beginner\venv\lib\site-
packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
[19]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	832
dense_1 (Dense)	(None, 32)	2,080
dense_2 (Dense)	(None, 1)	33

Total params: 2,945 (11.50 KB)

Trainable params: 2,945 (11.50 KB)

Non-trainable params: 0 (0.00 B)

```
[20]: opt = tf.keras.optimizers.Adam(learning_rate=0.01)
      loss = tf.keras.losses.BinaryCrossentropy()
```

```
[21]: # In order to forward and backward propagation we use to compile
      # compile the model

      model.compile(optimizer="adam", loss="binary_crossentropy",
                    metrics=['accuracy'])
```

```
[22]: ## Setup the Tensorboard
      from tensorflow.keras.callbacks import EarlyStopping, TensorBoard

      log_dir= "log/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
      tensorflow_callback=TensorBoard(log_dir=log_dir, histogram_freq=1)
```

Early Stopping: If we are train our model and let's say we are training it for 100 Epoch and after the training of 30 Epoch our loss value is not decreasing then no need to run the other 70 Epoch so for this we are using the EarlyStopping

```
[23]: ## setup Early Stopping
      early_stopping_callback = EarlyStopping(monitor='val_loss', patience=10,
                                             restore_best_weights=True)
```

2.1.1 Training the model

```
[24]: history = model.fit(
      X_train, y_train, validation_data=(X_test, y_test), epochs = 100,
      callbacks=[tensorflow_callback, early_stopping_callback]
      )
```

Epoch 1/100

250/250 4s 8ms/step -

accuracy: 0.8016 - loss: 0.4754 - val_accuracy: 0.8285 - val_loss: 0.3965

Epoch 2/100

250/250 2s 6ms/step -

accuracy: 0.8377 - loss: 0.3950 - val_accuracy: 0.8450 - val_loss: 0.3668

Epoch 3/100

250/250 3s 12ms/step -

accuracy: 0.8490 - loss: 0.3696 - val_accuracy: 0.8535 - val_loss: 0.3582

Epoch 4/100

250/250 5s 17ms/step -

accuracy: 0.8584 - loss: 0.3478 - val_accuracy: 0.8565 - val_loss: 0.3481
 Epoch 5/100
 250/250 2s 6ms/step -
 accuracy: 0.8576 - loss: 0.3460 - val_accuracy: 0.8580 - val_loss: 0.3496
 Epoch 6/100
 250/250 2s 6ms/step -
 accuracy: 0.8588 - loss: 0.3388 - val_accuracy: 0.8555 - val_loss: 0.3475
 Epoch 7/100
 250/250 2s 6ms/step -
 accuracy: 0.8619 - loss: 0.3401 - val_accuracy: 0.8630 - val_loss: 0.3423
 Epoch 8/100
 250/250 3s 9ms/step -
 accuracy: 0.8625 - loss: 0.3266 - val_accuracy: 0.8575 - val_loss: 0.3499
 Epoch 9/100
 250/250 1s 5ms/step -
 accuracy: 0.8563 - loss: 0.3330 - val_accuracy: 0.8610 - val_loss: 0.3403
 Epoch 10/100
 250/250 2s 8ms/step -
 accuracy: 0.8695 - loss: 0.3220 - val_accuracy: 0.8605 - val_loss: 0.3419
 Epoch 11/100
 250/250 4s 14ms/step -
 accuracy: 0.8675 - loss: 0.3190 - val_accuracy: 0.8535 - val_loss: 0.3488
 Epoch 12/100
 250/250 3s 7ms/step -
 accuracy: 0.8668 - loss: 0.3123 - val_accuracy: 0.8580 - val_loss: 0.3441
 Epoch 13/100
 250/250 2s 8ms/step -
 accuracy: 0.8643 - loss: 0.3242 - val_accuracy: 0.8630 - val_loss: 0.3420
 Epoch 14/100
 250/250 2s 6ms/step -
 accuracy: 0.8644 - loss: 0.3190 - val_accuracy: 0.8600 - val_loss: 0.3436
 Epoch 15/100
 250/250 2s 7ms/step -
 accuracy: 0.8724 - loss: 0.3119 - val_accuracy: 0.8595 - val_loss: 0.3411
 Epoch 16/100
 250/250 2s 8ms/step -
 accuracy: 0.8740 - loss: 0.3109 - val_accuracy: 0.8595 - val_loss: 0.3466
 Epoch 17/100
 250/250 2s 8ms/step -
 accuracy: 0.8741 - loss: 0.3101 - val_accuracy: 0.8585 - val_loss: 0.3427
 Epoch 18/100
 250/250 2s 7ms/step -
 accuracy: 0.8689 - loss: 0.3155 - val_accuracy: 0.8595 - val_loss: 0.3440
 Epoch 19/100
 250/250 2s 7ms/step -
 accuracy: 0.8719 - loss: 0.3048 - val_accuracy: 0.8590 - val_loss: 0.3459

```
[25]: model.save('model.h5')
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or  
`keras.saving.save_model(model)`. This file format is considered legacy. We  
recommend using instead the native Keras format, e.g.  
`model.save('my_model.keras')` or `keras.saving.save_model(model,  
'my_model.keras')`.
```

2.1.2 Loading Tensorflow Extension

```
[26]: %load_ext tensorboard
```

```
[28]: %tensorboard --logdir log/fit/20250406-205147
```

```
Reusing TensorBoard on port 6008 (pid 14308), started 0:00:16 ago. (Use '!kill_  
14308' to kill it.)
```

```
<IPython.core.display.HTML object>
```