

model

April 16, 2025

```
[1]: # Importing Required Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
```

```
[2]: # Load Dataset
data = pd.read_csv('data.csv')
data.head()
```

```
[2]:    id  x.radius_mean  x.texture_mean  x.perimeter_mean  x.area_mean  \
0    1         13.540         14.36         87.46         566.3
1    2         13.080         15.71         85.63         520.0
2    3          9.504         12.44         60.34         273.9
3    4         13.030         18.42         82.61         523.8
4    5          8.196         16.84         51.71         201.9

    x.smoothness_mean  x.compactness_mean  x.concavity_mean  \
0          0.09779          0.08129          0.06664
1          0.10750          0.12700          0.04568
2          0.10240          0.06492          0.02956
3          0.08983          0.03766          0.02562
4          0.08600          0.05943          0.01588

    x.concave_pts_mean  x.symmetry_mean  ...  x.texture_worst  \
0          0.047810          0.1885  ...          19.26
1          0.031100          0.1967  ...          20.49
2          0.020760          0.1815  ...          15.66
3          0.029230          0.1467  ...          22.81
4          0.005917          0.1769  ...          21.96

    x.perimeter_worst  x.area_worst  x.smoothness_worst  x.compactness_worst  \
0          99.70          711.2          0.14400          0.17730
```

1	96.09	630.5	0.13120	0.27760
2	65.13	314.9	0.13240	0.11480
3	84.46	545.9	0.09701	0.04619
4	57.26	242.2	0.12970	0.13570

	x.concavity_worst	x.concave_pts_worst	x.symmetry_worst	\
0	0.23900	0.12880	0.2977	
1	0.18900	0.07283	0.3184	
2	0.08867	0.06227	0.2450	
3	0.04833	0.05013	0.1987	
4	0.06880	0.02564	0.3105	

	x.fractal_dim_worst	diagnosis
0	0.07259	B
1	0.08183	B
2	0.07773	B
3	0.06169	B
4	0.07409	B

[5 rows x 32 columns]

```
[3]: # Drop Unnecessary Columns if present
if 'Unnamed: 32' in data.columns:
    data.drop('Unnamed: 32', axis=1, inplace=True)
if 'id' in data.columns:
    data.drop('id', axis=1, inplace=True)
```

```
[4]: # Encode Labels
data['diagnosis'] = data['diagnosis'].map({'M': 1, 'B': 0})
data['diagnosis']
```

```
[4]: 0      0
      1      0
      2      0
      3      0
      4      0
      ..
     564     1
     565     1
     566     1
     567     1
     568     1
      Name: diagnosis, Length: 569, dtype: int64
```

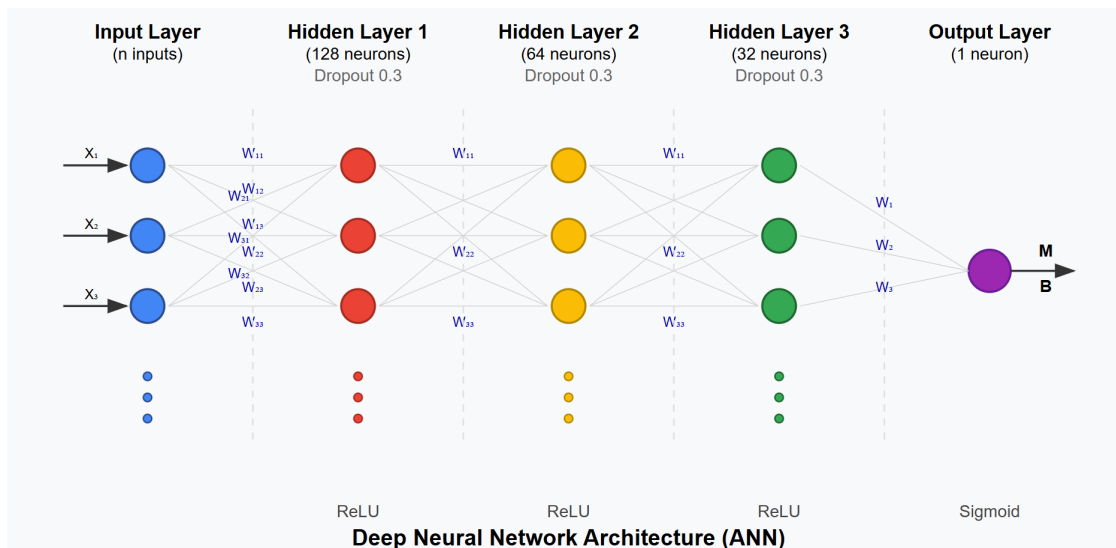
```
[5]: # Split features and labels
X = data.drop('diagnosis', axis=1)
y = data['diagnosis']
```

```
# Standardize the Data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
random_state=42)
```

```
[6]: # Build Deep Learning Model
model = Sequential([
    Dense(128, activation='relu', input_shape=(X_train.
shape[1],), name='Input_layer'),
    Dropout(0.3),
    Dense(64, activation='relu', name='First_hidden_layer'),
    Dropout(0.3),
    Dense(32, activation='relu', name='Second_hidden_layer'),
    Dropout(0.3),
    Dense(1, activation='sigmoid', name='Output_layer')
])
```

c:\Users\Uditya\Desktop\Deep Learning\Deep Learning for Beginner\venv\lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)



```
[7]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
Input_layer (Dense)	(None, 128)	3,968
dropout (Dropout)	(None, 128)	0
First_hidden_layer (Dense)	(None, 64)	8,256
dropout_1 (Dropout)	(None, 64)	0
Second_hidden_layer (Dense)	(None, 32)	2,080
dropout_2 (Dropout)	(None, 32)	0
Output_layer (Dense)	(None, 1)	33

Total params: 14,337 (56.00 KB)

Trainable params: 14,337 (56.00 KB)

Non-trainable params: 0 (0.00 B)

```
[8]: # Compile Model
model.compile(optimizer='adam', loss='binary_crossentropy',
             metrics=['accuracy'])

# Train Model
history = model.fit(X_train, y_train, epochs=100, batch_size=32,
                  validation_split=0.2, verbose=1)

# Evaluate Model
y_pred = (model.predict(X_test) > 0.5).astype("int32")
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

```
Epoch 1/100
12/12          5s 63ms/step -
accuracy: 0.7209 - loss: 0.5825 - val_accuracy: 0.9341 - val_loss: 0.3944
Epoch 2/100
12/12          1s 22ms/step -
accuracy: 0.8687 - loss: 0.3941 - val_accuracy: 0.9451 - val_loss: 0.2371
Epoch 3/100
12/12          0s 16ms/step -
```

accuracy: 0.9257 - loss: 0.2348 - val_accuracy: 0.9451 - val_loss: 0.1562
 Epoch 4/100
 12/12 0s 26ms/step -
 accuracy: 0.9498 - loss: 0.1543 - val_accuracy: 0.9451 - val_loss: 0.1163
 Epoch 5/100
 12/12 0s 26ms/step -
 accuracy: 0.9728 - loss: 0.1069 - val_accuracy: 0.9670 - val_loss: 0.0902
 Epoch 6/100
 12/12 0s 33ms/step -
 accuracy: 0.9675 - loss: 0.0967 - val_accuracy: 0.9670 - val_loss: 0.0793
 Epoch 7/100
 12/12 1s 22ms/step -
 accuracy: 0.9610 - loss: 0.1094 - val_accuracy: 0.9670 - val_loss: 0.0765
 Epoch 8/100
 12/12 0s 16ms/step -
 accuracy: 0.9777 - loss: 0.0987 - val_accuracy: 0.9780 - val_loss: 0.0875
 Epoch 9/100
 12/12 0s 16ms/step -
 accuracy: 0.9786 - loss: 0.0690 - val_accuracy: 0.9780 - val_loss: 0.0913
 Epoch 10/100
 12/12 0s 14ms/step -
 accuracy: 0.9889 - loss: 0.0575 - val_accuracy: 0.9780 - val_loss: 0.0935
 Epoch 11/100
 12/12 0s 15ms/step -
 accuracy: 0.9846 - loss: 0.0554 - val_accuracy: 0.9780 - val_loss: 0.0967
 Epoch 12/100
 12/12 0s 25ms/step -
 accuracy: 0.9901 - loss: 0.0487 - val_accuracy: 0.9780 - val_loss: 0.0929
 Epoch 13/100
 12/12 0s 15ms/step -
 accuracy: 0.9903 - loss: 0.0336 - val_accuracy: 0.9780 - val_loss: 0.0965
 Epoch 14/100
 12/12 0s 27ms/step -
 accuracy: 0.9919 - loss: 0.0410 - val_accuracy: 0.9780 - val_loss: 0.0961
 Epoch 15/100
 12/12 0s 18ms/step -
 accuracy: 0.9878 - loss: 0.0435 - val_accuracy: 0.9780 - val_loss: 0.0986
 Epoch 16/100
 12/12 0s 20ms/step -
 accuracy: 0.9933 - loss: 0.0407 - val_accuracy: 0.9780 - val_loss: 0.0956
 Epoch 17/100
 12/12 0s 17ms/step -
 accuracy: 0.9915 - loss: 0.0389 - val_accuracy: 0.9780 - val_loss: 0.1009
 Epoch 18/100
 12/12 0s 22ms/step -
 accuracy: 0.9888 - loss: 0.0477 - val_accuracy: 0.9780 - val_loss: 0.1136
 Epoch 19/100
 12/12 0s 18ms/step -

accuracy: 0.9926 - loss: 0.0411 - val_accuracy: 0.9780 - val_loss: 0.1106
 Epoch 20/100
 12/12 0s 17ms/step -
 accuracy: 0.9954 - loss: 0.0303 - val_accuracy: 0.9780 - val_loss: 0.1085
 Epoch 21/100
 12/12 0s 21ms/step -
 accuracy: 0.9876 - loss: 0.0295 - val_accuracy: 0.9780 - val_loss: 0.1118
 Epoch 22/100
 12/12 0s 24ms/step -
 accuracy: 0.9969 - loss: 0.0225 - val_accuracy: 0.9780 - val_loss: 0.1168
 Epoch 23/100
 12/12 0s 18ms/step -
 accuracy: 0.9977 - loss: 0.0173 - val_accuracy: 0.9780 - val_loss: 0.1120
 Epoch 24/100
 12/12 0s 18ms/step -
 accuracy: 0.9914 - loss: 0.0251 - val_accuracy: 0.9780 - val_loss: 0.1063
 Epoch 25/100
 12/12 0s 17ms/step -
 accuracy: 0.9932 - loss: 0.0283 - val_accuracy: 0.9780 - val_loss: 0.1122
 Epoch 26/100
 12/12 0s 18ms/step -
 accuracy: 0.9953 - loss: 0.0236 - val_accuracy: 0.9780 - val_loss: 0.1180
 Epoch 27/100
 12/12 0s 19ms/step -
 accuracy: 0.9920 - loss: 0.0226 - val_accuracy: 0.9670 - val_loss: 0.1214
 Epoch 28/100
 12/12 0s 19ms/step -
 accuracy: 0.9973 - loss: 0.0181 - val_accuracy: 0.9780 - val_loss: 0.1203
 Epoch 29/100
 12/12 1s 20ms/step -
 accuracy: 0.9884 - loss: 0.0260 - val_accuracy: 0.9780 - val_loss: 0.1311
 Epoch 30/100
 12/12 0s 18ms/step -
 accuracy: 0.9959 - loss: 0.0163 - val_accuracy: 0.9780 - val_loss: 0.1375
 Epoch 31/100
 12/12 0s 29ms/step -
 accuracy: 0.9978 - loss: 0.0102 - val_accuracy: 0.9780 - val_loss: 0.1430
 Epoch 32/100
 12/12 1s 25ms/step -
 accuracy: 0.9901 - loss: 0.0378 - val_accuracy: 0.9780 - val_loss: 0.1400
 Epoch 33/100
 12/12 0s 20ms/step -
 accuracy: 0.9909 - loss: 0.0201 - val_accuracy: 0.9780 - val_loss: 0.1355
 Epoch 34/100
 12/12 0s 18ms/step -
 accuracy: 0.9969 - loss: 0.0143 - val_accuracy: 0.9780 - val_loss: 0.1480
 Epoch 35/100
 12/12 0s 19ms/step -

accuracy: 0.9901 - loss: 0.0272 - val_accuracy: 0.9780 - val_loss: 0.1484
 Epoch 36/100
 12/12 0s 20ms/step -
 accuracy: 1.0000 - loss: 0.0082 - val_accuracy: 0.9780 - val_loss: 0.1515
 Epoch 37/100
 12/12 0s 20ms/step -
 accuracy: 0.9985 - loss: 0.0051 - val_accuracy: 0.9780 - val_loss: 0.1591
 Epoch 38/100
 12/12 0s 22ms/step -
 accuracy: 0.9963 - loss: 0.0101 - val_accuracy: 0.9780 - val_loss: 0.1561
 Epoch 39/100
 12/12 0s 33ms/step -
 accuracy: 1.0000 - loss: 0.0085 - val_accuracy: 0.9670 - val_loss: 0.1586
 Epoch 40/100
 12/12 0s 21ms/step -
 accuracy: 0.9969 - loss: 0.0135 - val_accuracy: 0.9780 - val_loss: 0.1539
 Epoch 41/100
 12/12 0s 18ms/step -
 accuracy: 0.9947 - loss: 0.0172 - val_accuracy: 0.9780 - val_loss: 0.1639
 Epoch 42/100
 12/12 0s 18ms/step -
 accuracy: 0.9994 - loss: 0.0046 - val_accuracy: 0.9780 - val_loss: 0.1739
 Epoch 43/100
 12/12 0s 16ms/step -
 accuracy: 1.0000 - loss: 0.0043 - val_accuracy: 0.9780 - val_loss: 0.1751
 Epoch 44/100
 12/12 0s 17ms/step -
 accuracy: 1.0000 - loss: 0.0032 - val_accuracy: 0.9780 - val_loss: 0.1796
 Epoch 45/100
 12/12 0s 17ms/step -
 accuracy: 1.0000 - loss: 0.0037 - val_accuracy: 0.9780 - val_loss: 0.1842
 Epoch 46/100
 12/12 0s 20ms/step -
 accuracy: 1.0000 - loss: 0.0026 - val_accuracy: 0.9780 - val_loss: 0.1892
 Epoch 47/100
 12/12 0s 20ms/step -
 accuracy: 0.9973 - loss: 0.0059 - val_accuracy: 0.9780 - val_loss: 0.1892
 Epoch 48/100
 12/12 0s 21ms/step -
 accuracy: 1.0000 - loss: 0.0046 - val_accuracy: 0.9670 - val_loss: 0.1890
 Epoch 49/100
 12/12 0s 21ms/step -
 accuracy: 1.0000 - loss: 0.0040 - val_accuracy: 0.9780 - val_loss: 0.1832
 Epoch 50/100
 12/12 0s 20ms/step -
 accuracy: 0.9996 - loss: 0.0044 - val_accuracy: 0.9670 - val_loss: 0.2014
 Epoch 51/100
 12/12 0s 18ms/step -

```

accuracy: 1.0000 - loss: 0.0035 - val_accuracy: 0.9670 - val_loss: 0.2209
Epoch 52/100
12/12          0s 19ms/step -
accuracy: 1.0000 - loss: 0.0031 - val_accuracy: 0.9670 - val_loss: 0.2415
Epoch 53/100
12/12          0s 23ms/step -
accuracy: 1.0000 - loss: 0.0027 - val_accuracy: 0.9670 - val_loss: 0.2809
Epoch 54/100
12/12          0s 33ms/step -
accuracy: 1.0000 - loss: 0.0028 - val_accuracy: 0.9670 - val_loss: 0.3221
Epoch 55/100
12/12          1s 25ms/step -
accuracy: 0.9897 - loss: 0.0289 - val_accuracy: 0.9780 - val_loss: 0.1959
Epoch 56/100
12/12          1s 20ms/step -
accuracy: 0.9923 - loss: 0.0148 - val_accuracy: 0.9670 - val_loss: 0.2122
Epoch 57/100
12/12          0s 20ms/step -
accuracy: 1.0000 - loss: 0.0070 - val_accuracy: 0.9670 - val_loss: 0.2424
Epoch 58/100
12/12          0s 19ms/step -
accuracy: 1.0000 - loss: 0.0022 - val_accuracy: 0.9780 - val_loss: 0.2228
Epoch 59/100
12/12          0s 21ms/step -
accuracy: 0.9972 - loss: 0.0086 - val_accuracy: 0.9780 - val_loss: 0.2187
Epoch 60/100
12/12          0s 20ms/step -
accuracy: 1.0000 - loss: 0.0022 - val_accuracy: 0.9780 - val_loss: 0.2235
Epoch 61/100
12/12          0s 21ms/step -
accuracy: 1.0000 - loss: 0.0012 - val_accuracy: 0.9780 - val_loss: 0.2323
Epoch 62/100
12/12          0s 21ms/step -
accuracy: 0.9947 - loss: 0.0117 - val_accuracy: 0.9670 - val_loss: 0.3633
Epoch 63/100
12/12          0s 21ms/step -
accuracy: 0.9994 - loss: 0.0042 - val_accuracy: 0.9670 - val_loss: 0.3622
Epoch 64/100
12/12          0s 33ms/step -
accuracy: 1.0000 - loss: 0.0037 - val_accuracy: 0.9780 - val_loss: 0.2298
Epoch 65/100
12/12          1s 33ms/step -
accuracy: 1.0000 - loss: 8.6143e-04 - val_accuracy: 0.9780 - val_loss: 0.2147
Epoch 66/100
12/12          1s 31ms/step -
accuracy: 1.0000 - loss: 0.0056 - val_accuracy: 0.9780 - val_loss: 0.2166
Epoch 67/100
12/12          1s 36ms/step -

```


accuracy: 0.9973 - loss: 0.0131 - val_accuracy: 0.9670 - val_loss: 0.3318
 Epoch 68/100
 12/12 1s 29ms/step -
 accuracy: 0.9973 - loss: 0.0039 - val_accuracy: 0.9560 - val_loss: 0.3811
 Epoch 69/100
 12/12 0s 23ms/step -
 accuracy: 0.9947 - loss: 0.0078 - val_accuracy: 0.9560 - val_loss: 0.3952
 Epoch 70/100
 12/12 0s 36ms/step -
 accuracy: 1.0000 - loss: 0.0029 - val_accuracy: 0.9560 - val_loss: 0.3964
 Epoch 71/100
 12/12 1s 23ms/step -
 accuracy: 1.0000 - loss: 0.0018 - val_accuracy: 0.9560 - val_loss: 0.3896
 Epoch 72/100
 12/12 0s 20ms/step -
 accuracy: 1.0000 - loss: 0.0045 - val_accuracy: 0.9560 - val_loss: 0.4393
 Epoch 73/100
 12/12 0s 16ms/step -
 accuracy: 1.0000 - loss: 0.0015 - val_accuracy: 0.9560 - val_loss: 0.4690
 Epoch 74/100
 12/12 0s 16ms/step -
 accuracy: 0.9978 - loss: 0.0080 - val_accuracy: 0.9670 - val_loss: 0.2628
 Epoch 75/100
 12/12 0s 19ms/step -
 accuracy: 0.9947 - loss: 0.0113 - val_accuracy: 0.9670 - val_loss: 0.2435
 Epoch 76/100
 12/12 0s 19ms/step -
 accuracy: 1.0000 - loss: 0.0026 - val_accuracy: 0.9670 - val_loss: 0.2326
 Epoch 77/100
 12/12 0s 22ms/step -
 accuracy: 0.9982 - loss: 0.0052 - val_accuracy: 0.9780 - val_loss: 0.2300
 Epoch 78/100
 12/12 0s 19ms/step -
 accuracy: 0.9986 - loss: 0.0038 - val_accuracy: 0.9780 - val_loss: 0.2258
 Epoch 79/100
 12/12 0s 18ms/step -
 accuracy: 1.0000 - loss: 0.0017 - val_accuracy: 0.9780 - val_loss: 0.2214
 Epoch 80/100
 12/12 1s 21ms/step -
 accuracy: 1.0000 - loss: 0.0055 - val_accuracy: 0.9780 - val_loss: 0.1910
 Epoch 81/100
 12/12 0s 18ms/step -
 accuracy: 1.0000 - loss: 0.0019 - val_accuracy: 0.9780 - val_loss: 0.1781
 Epoch 82/100
 12/12 0s 24ms/step -
 accuracy: 1.0000 - loss: 0.0017 - val_accuracy: 0.9780 - val_loss: 0.1892
 Epoch 83/100
 12/12 0s 21ms/step -

accuracy: 1.0000 - loss: 0.0022 - val_accuracy: 0.9780 - val_loss: 0.1755
 Epoch 84/100
 12/12 0s 18ms/step -
 accuracy: 1.0000 - loss: 0.0019 - val_accuracy: 0.9780 - val_loss: 0.1749
 Epoch 85/100
 12/12 0s 17ms/step -
 accuracy: 0.9923 - loss: 0.0116 - val_accuracy: 0.9780 - val_loss: 0.1824
 Epoch 86/100
 12/12 0s 18ms/step -
 accuracy: 1.0000 - loss: 0.0012 - val_accuracy: 0.9780 - val_loss: 0.1858
 Epoch 87/100
 12/12 0s 28ms/step -
 accuracy: 1.0000 - loss: 0.0010 - val_accuracy: 0.9780 - val_loss: 0.1908
 Epoch 88/100
 12/12 1s 23ms/step -
 accuracy: 1.0000 - loss: 0.0033 - val_accuracy: 0.9780 - val_loss: 0.2192
 Epoch 89/100
 12/12 0s 21ms/step -
 accuracy: 0.9973 - loss: 0.0036 - val_accuracy: 0.9780 - val_loss: 0.2307
 Epoch 90/100
 12/12 0s 18ms/step -
 accuracy: 1.0000 - loss: 0.0032 - val_accuracy: 0.9780 - val_loss: 0.2390
 Epoch 91/100
 12/12 0s 20ms/step -
 accuracy: 1.0000 - loss: 0.0052 - val_accuracy: 0.9780 - val_loss: 0.2384
 Epoch 92/100
 12/12 0s 19ms/step -
 accuracy: 1.0000 - loss: 6.4148e-04 - val_accuracy: 0.9780 - val_loss: 0.2426
 Epoch 93/100
 12/12 0s 29ms/step -
 accuracy: 1.0000 - loss: 3.3552e-04 - val_accuracy: 0.9780 - val_loss: 0.2467
 Epoch 94/100
 12/12 0s 20ms/step -
 accuracy: 0.9947 - loss: 0.0062 - val_accuracy: 0.9670 - val_loss: 0.3416
 Epoch 95/100
 12/12 0s 20ms/step -
 accuracy: 1.0000 - loss: 2.9441e-04 - val_accuracy: 0.9670 - val_loss: 0.3676
 Epoch 96/100
 12/12 0s 20ms/step -
 accuracy: 1.0000 - loss: 0.0027 - val_accuracy: 0.9670 - val_loss: 0.3203
 Epoch 97/100
 12/12 0s 18ms/step -
 accuracy: 1.0000 - loss: 0.0033 - val_accuracy: 0.9560 - val_loss: 0.3126
 Epoch 98/100
 12/12 0s 20ms/step -
 accuracy: 1.0000 - loss: 9.1326e-04 - val_accuracy: 0.9560 - val_loss: 0.3131
 Epoch 99/100
 12/12 0s 20ms/step -

```
accuracy: 1.0000 - loss: 0.0017 - val_accuracy: 0.9560 - val_loss: 0.3290
Epoch 100/100
12/12          0s 17ms/step -
accuracy: 0.9982 - loss: 0.0038 - val_accuracy: 0.9560 - val_loss: 0.3404
4/4           0s 46ms/step
```

Confusion Matrix:

```
[[71  0]
 [ 3 40]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.96	1.00	0.98	71
1	1.00	0.93	0.96	43
accuracy			0.97	114
macro avg	0.98	0.97	0.97	114
weighted avg	0.97	0.97	0.97	114

```
[9]: import os

# Create results folder if not exists
os.makedirs("results", exist_ok=True)


# Plot Accuracy and Loss
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Val Accuracy')
plt.title('Model Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('Model Loss')
plt.legend()

plt.tight_layout()


# Save the plot
plt.savefig("results/training_metrics.png")
plt.show()
```



```
[10]: # Plotting only Loss (Gradient Descent Visualization)
import os
os.makedirs("results", exist_ok=True)

plt.figure(figsize=(6, 4))
plt.plot(history.history['loss'], label='Train Loss', color='blue')
plt.plot(history.history['val_loss'], label='Validation Loss', color='orange')
plt.title('Loss Curve (Gradient Descent)')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)

plt.savefig("results/loss_gradient_descent.png")
plt.show()
```



```
[12]: # Prepare Dataset
train_dataset = tf.data.Dataset.from_tensor_slices((X_train, y_train)).batch(16)

# Optimizer & Loss
loss_fn = tf.keras.losses.BinaryCrossentropy()
optimizer = tf.keras.optimizers.Adam()

# For Tracking
gradient_norms = {layer.name: [] for layer in model.layers if len(layer.
    ↪ trainable_variables) > 0}
losses = []

# Training Loop
epochs = 100
for epoch in range(epochs):
    print(f"\nEpoch {epoch+1}/{epochs}")
    epoch_losses = []

    for step, (x_batch, y_batch) in enumerate(train_dataset):
        with tf.GradientTape() as tape:
            logits = model(x_batch, training=True)
            loss_value = loss_fn(y_batch, logits)
```

```

grads = tape.gradient(loss_value, model.trainable_variables)
optimizer.apply_gradients(zip(grads, model.trainable_variables))
epoch_losses.append(loss_value.numpy())

# Save gradient norms
idx = 0
for layer in model.layers:
    if len(layer.trainable_variables) > 0:
        for var in layer.trainable_variables:
            grad = grads[idx]
            norm = tf.norm(grad).numpy() if grad is not None else 0
            gradient_norms[layer.name].append(norm)
            idx += 1

# Average loss of epoch
epoch_loss = np.mean(epoch_losses)
losses.append(epoch_loss)
print(f"Loss: {epoch_loss:.4f}")

# Plotting
plt.figure(figsize=(14, 5), dpi=200)

# Loss vs Epoch
plt.subplot(1, 2, 1)
plt.plot(range(1, epochs+1), losses, marker='o', color='blue')
plt.title("Loss vs Epochs")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.grid(True)

# Gradient Norm vs Epoch (per layer)
plt.subplot(1, 2, 2)
steps_per_epoch = len(train_dataset)

for layer_name, norms in gradient_norms.items():
    norms = np.array(norms)
    avg_per_epoch = [np.mean(norms[i*steps_per_epoch:(i+1)*steps_per_epoch])]
    for i in range(epochs)]
    plt.plot(range(1, epochs+1), avg_per_epoch, label=layer_name)

plt.title("Gradient Norm vs Epochs (per layer)")
plt.xlabel("Epochs")
plt.ylabel("Avg Gradient Norm")
plt.legend()
plt.grid(True)

plt.tight_layout()

```

```
plt.savefig("results/gradient_norm_per_epoch.png")  
plt.show()
```

Epoch 1/100
Loss: 0.0208

Epoch 2/100
Loss: 0.0104

Epoch 3/100
Loss: 0.0570

Epoch 4/100
Loss: 0.0029

Epoch 5/100
Loss: 0.0206

Epoch 6/100
Loss: 0.0036

Epoch 7/100
Loss: 0.0037

Epoch 8/100
Loss: 0.0080

Epoch 9/100
Loss: 0.0074

Epoch 10/100
Loss: 0.0125

Epoch 11/100
Loss: 0.0087

Epoch 12/100
Loss: 0.0068

Epoch 13/100
Loss: 0.0076

Epoch 14/100
Loss: 0.0115

Epoch 15/100
Loss: 0.0447

Epoch 16/100
Loss: 0.0202

Epoch 17/100
Loss: 0.0035

Epoch 18/100
Loss: 0.0025

Epoch 19/100
Loss: 0.0119

Epoch 20/100
Loss: 0.0033

Epoch 21/100
Loss: 0.0061

Epoch 22/100
Loss: 0.0084

Epoch 23/100
Loss: 0.0067

Epoch 24/100
Loss: 0.0059

Epoch 25/100
Loss: 0.0030

Epoch 26/100
Loss: 0.0028

Epoch 27/100
Loss: 0.0020

Epoch 28/100
Loss: 0.0018

Epoch 29/100
Loss: 0.0007

Epoch 30/100
Loss: 0.0022

Epoch 31/100
Loss: 0.0013

Epoch 32/100
Loss: 0.0015

Epoch 33/100
Loss: 0.0007

Epoch 34/100
Loss: 0.0009

Epoch 35/100
Loss: 0.0011

Epoch 36/100
Loss: 0.0027

Epoch 37/100
Loss: 0.0012

Epoch 38/100
Loss: 0.0006

Epoch 39/100
Loss: 0.0014

Epoch 40/100
Loss: 0.0014

Epoch 41/100
Loss: 0.0003

Epoch 42/100
Loss: 0.0003

Epoch 43/100
Loss: 0.0003

Epoch 44/100
Loss: 0.0003

Epoch 45/100
Loss: 0.0009

Epoch 46/100
Loss: 0.0043

Epoch 47/100
Loss: 0.0029

Epoch 48/100
Loss: 0.0142

Epoch 49/100
Loss: 0.0807

Epoch 50/100
Loss: 0.0080

Epoch 51/100
Loss: 0.0134

Epoch 52/100
Loss: 0.0056

Epoch 53/100
Loss: 0.0121

Epoch 54/100
Loss: 0.0021

Epoch 55/100
Loss: 0.0023

Epoch 56/100
Loss: 0.0013

Epoch 57/100
Loss: 0.0007

Epoch 58/100
Loss: 0.0010

Epoch 59/100
Loss: 0.0015

Epoch 60/100
Loss: 0.0108

Epoch 61/100
Loss: 0.0009

Epoch 62/100
Loss: 0.0006

Epoch 63/100
Loss: 0.0011

Epoch 64/100
Loss: 0.0112

Epoch 65/100
Loss: 0.0016

Epoch 66/100
Loss: 0.0017

Epoch 67/100
Loss: 0.0008

Epoch 68/100
Loss: 0.0019

Epoch 69/100
Loss: 0.0009

Epoch 70/100
Loss: 0.0006

Epoch 71/100
Loss: 0.0014

Epoch 72/100
Loss: 0.0012

Epoch 73/100
Loss: 0.0007

Epoch 74/100
Loss: 0.0002

Epoch 75/100
Loss: 0.0133

Epoch 76/100
Loss: 0.0053

Epoch 77/100
Loss: 0.0246

Epoch 78/100
Loss: 0.0059

Epoch 79/100
Loss: 0.0012

Epoch 80/100
Loss: 0.0024

Epoch 81/100
Loss: 0.0005

Epoch 82/100
Loss: 0.0012

Epoch 83/100
Loss: 0.0011

Epoch 84/100
Loss: 0.0009

Epoch 85/100
Loss: 0.0004

Epoch 86/100
Loss: 0.0004

Epoch 87/100
Loss: 0.0004

Epoch 88/100
Loss: 0.0005

Epoch 89/100
Loss: 0.0001

Epoch 90/100
Loss: 0.0006

Epoch 91/100
Loss: 0.0002

Epoch 92/100
Loss: 0.0014

Epoch 93/100
Loss: 0.0008

Epoch 94/100
Loss: 0.0008

Epoch 95/100
Loss: 0.0002

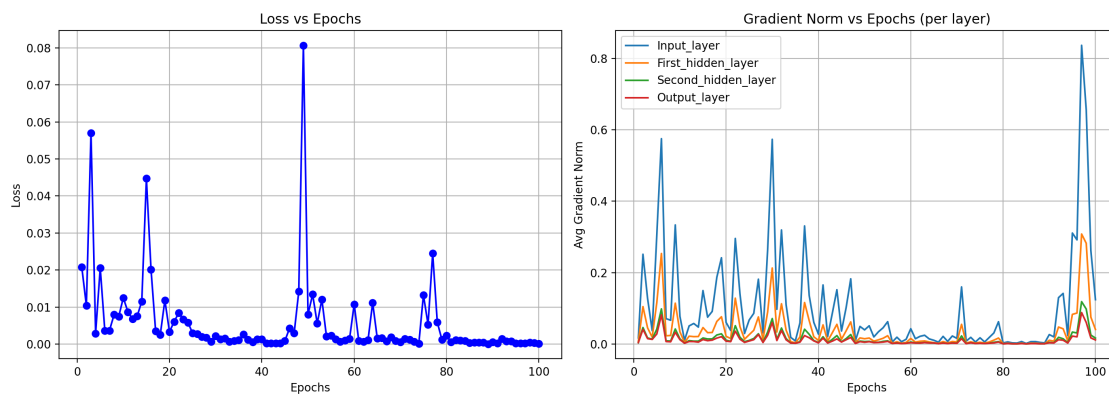
Epoch 96/100
Loss: 0.0003

Epoch 97/100
Loss: 0.0002

Epoch 98/100
Loss: 0.0005

Epoch 99/100
Loss: 0.0004


Epoch 100/100
Loss: 0.0002



```
[13]: # Create result directory if not exists
import os
os.makedirs("results", exist_ok=True)

# Plot gradient norms for each layer
plt.figure(figsize=(10, 6), dpi= 200)
for layer_name, norms in gradient_norms.items():
    plt.plot(norms, label=layer_name)

plt.title("Gradient Norms per Layer across Batches")
plt.xlabel("Batch Steps")
plt.ylabel("Gradient Norm")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig("results/gradient_flow_per_layer.png")
plt.show()
```



0.1 Binary Cross Entropy Loss

core of **gradient descent** update:

$$W_{t+1} = W_t - \eta \cdot \frac{\partial L}{\partial W}$$

core of **gradient descent** update:

$$W_{t+1} = W_t - \eta \cdot \frac{\partial L}{\partial W}$$

core of gradient descent update:

$$W_{t+1} = W_t - \eta \cdot \frac{\partial L}{\partial W}$$

```
[14]: loss_fn = tf.keras.losses.BinaryCrossentropy()
optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)

# To store gradient norms
gradient_norms = []

# Convert data to tensors
X_tensor = tf.convert_to_tensor(X_train, dtype=tf.float32)
y_tensor = tf.convert_to_tensor(y_train.values.reshape(-1, 1), dtype=tf.float32)

# Training Loop for 100 epochs
for epoch in range(100):
    with tf.GradientTape() as tape:
        predictions = model(X_tensor, training=True)
        loss_value = loss_fn(y_tensor, predictions)

    # Compute gradients
    grads = tape.gradient(loss_value, model.trainable_weights)

    # Record gradient norms
    grad_norm = np.mean([tf.norm(g).numpy() for g in grads if g is not None])
    gradient_norms.append(grad_norm)

    # Apply gradients
    optimizer.apply_gradients(zip(grads, model.trainable_weights))

    print(f"Epoch {epoch+1}: Loss = {loss_value:.4f}, Avg. Grad Norm = {grad_norm:.4f}")

# Plot gradient norms
plt.figure(figsize=(12, 7), dpi=250)
plt.plot(gradient_norms, marker='o', color='green')
plt.title("Average Gradient Norm per Epoch")
plt.xlabel("Epoch")
plt.ylabel("Gradient Norm")
plt.grid(True)
plt.tight_layout()
plt.savefig("results/gradient_flow.png")
```

```
plt.show()
```

```
Epoch 1: Loss = 0.0001, Avg. Grad Norm = 0.0003
Epoch 2: Loss = 0.0001, Avg. Grad Norm = 0.0004
Epoch 3: Loss = 0.0004, Avg. Grad Norm = 0.0041
Epoch 4: Loss = 0.0002, Avg. Grad Norm = 0.0006
Epoch 5: Loss = 0.0001, Avg. Grad Norm = 0.0015
Epoch 6: Loss = 0.0001, Avg. Grad Norm = 0.0005
Epoch 7: Loss = 0.0001, Avg. Grad Norm = 0.0010
Epoch 8: Loss = 0.0007, Avg. Grad Norm = 0.0078
Epoch 9: Loss = 0.0008, Avg. Grad Norm = 0.0062
Epoch 10: Loss = 0.0001, Avg. Grad Norm = 0.0003
Epoch 11: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 12: Loss = 0.0003, Avg. Grad Norm = 0.0031
Epoch 13: Loss = 0.0006, Avg. Grad Norm = 0.0046
Epoch 14: Loss = 0.0001, Avg. Grad Norm = 0.0003
Epoch 15: Loss = 0.0008, Avg. Grad Norm = 0.0081
Epoch 16: Loss = 0.0026, Avg. Grad Norm = 0.0138
Epoch 17: Loss = 0.0001, Avg. Grad Norm = 0.0010
Epoch 18: Loss = 0.0003, Avg. Grad Norm = 0.0024
Epoch 19: Loss = 0.0001, Avg. Grad Norm = 0.0005
Epoch 20: Loss = 0.0002, Avg. Grad Norm = 0.0012
Epoch 21: Loss = 0.0060, Avg. Grad Norm = 0.0569
Epoch 22: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 23: Loss = 0.0018, Avg. Grad Norm = 0.0229
Epoch 24: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 25: Loss = 0.0002, Avg. Grad Norm = 0.0009
Epoch 26: Loss = 0.0001, Avg. Grad Norm = 0.0006
Epoch 27: Loss = 0.0001, Avg. Grad Norm = 0.0003
Epoch 28: Loss = 0.0001, Avg. Grad Norm = 0.0006
Epoch 29: Loss = 0.0001, Avg. Grad Norm = 0.0016
Epoch 30: Loss = 0.0001, Avg. Grad Norm = 0.0005
Epoch 31: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 32: Loss = 0.0004, Avg. Grad Norm = 0.0040
Epoch 33: Loss = 0.0001, Avg. Grad Norm = 0.0005
Epoch 34: Loss = 0.0002, Avg. Grad Norm = 0.0009
Epoch 35: Loss = 0.0003, Avg. Grad Norm = 0.0027
Epoch 36: Loss = 0.0001, Avg. Grad Norm = 0.0003
Epoch 37: Loss = 0.0005, Avg. Grad Norm = 0.0068
Epoch 38: Loss = 0.0001, Avg. Grad Norm = 0.0012
Epoch 39: Loss = 0.0001, Avg. Grad Norm = 0.0004
Epoch 40: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 41: Loss = 0.0001, Avg. Grad Norm = 0.0003
Epoch 42: Loss = 0.0002, Avg. Grad Norm = 0.0005
Epoch 43: Loss = 0.0002, Avg. Grad Norm = 0.0025
Epoch 44: Loss = 0.0001, Avg. Grad Norm = 0.0007
Epoch 45: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 46: Loss = 0.0000, Avg. Grad Norm = 0.0004
```


Epoch 47: Loss = 0.0010, Avg. Grad Norm = 0.0360
Epoch 48: Loss = 0.0001, Avg. Grad Norm = 0.0005
Epoch 49: Loss = 0.0001, Avg. Grad Norm = 0.0003
Epoch 50: Loss = 0.0000, Avg. Grad Norm = 0.0004
Epoch 51: Loss = 0.0002, Avg. Grad Norm = 0.0025
Epoch 52: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 53: Loss = 0.0003, Avg. Grad Norm = 0.0019
Epoch 54: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 55: Loss = 0.0000, Avg. Grad Norm = 0.0003
Epoch 56: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 57: Loss = 0.0001, Avg. Grad Norm = 0.0005
Epoch 58: Loss = 0.0146, Avg. Grad Norm = 0.0521
Epoch 59: Loss = 0.0011, Avg. Grad Norm = 0.0094
Epoch 60: Loss = 0.0001, Avg. Grad Norm = 0.0007
Epoch 61: Loss = 0.0023, Avg. Grad Norm = 0.0351
Epoch 62: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 63: Loss = 0.0003, Avg. Grad Norm = 0.0011
Epoch 64: Loss = 0.0001, Avg. Grad Norm = 0.0008
Epoch 65: Loss = 0.0000, Avg. Grad Norm = 0.0005
Epoch 66: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 67: Loss = 0.0003, Avg. Grad Norm = 0.0034
Epoch 68: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 69: Loss = 0.0001, Avg. Grad Norm = 0.0003
Epoch 70: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 71: Loss = 0.0001, Avg. Grad Norm = 0.0004
Epoch 72: Loss = 0.0003, Avg. Grad Norm = 0.0037
Epoch 73: Loss = 0.0001, Avg. Grad Norm = 0.0014
Epoch 74: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 75: Loss = 0.0001, Avg. Grad Norm = 0.0019
Epoch 76: Loss = 0.0002, Avg. Grad Norm = 0.0015
Epoch 77: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 78: Loss = 0.0002, Avg. Grad Norm = 0.0011
Epoch 79: Loss = 0.0001, Avg. Grad Norm = 0.0006
Epoch 80: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 81: Loss = 0.0001, Avg. Grad Norm = 0.0002
Epoch 82: Loss = 0.0001, Avg. Grad Norm = 0.0009
Epoch 83: Loss = 0.0001, Avg. Grad Norm = 0.0009
Epoch 84: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 85: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 86: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 87: Loss = 0.0000, Avg. Grad Norm = 0.0005
Epoch 88: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 89: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 90: Loss = 0.0038, Avg. Grad Norm = 0.0328
Epoch 91: Loss = 0.0009, Avg. Grad Norm = 0.0079
Epoch 92: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 93: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 94: Loss = 0.0000, Avg. Grad Norm = 0.0002

Epoch 95: Loss = 0.0003, Avg. Grad Norm = 0.0029
Epoch 96: Loss = 0.0001, Avg. Grad Norm = 0.0003
Epoch 97: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 98: Loss = 0.0001, Avg. Grad Norm = 0.0005
Epoch 99: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 100: Loss = 0.0001, Avg. Grad Norm = 0.0011

