

model

April 17, 2025

```
[1]: # Importing Required Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
```

```
[2]: # Load Dataset
data = pd.read_csv('data.csv')
data.head()
```

```
[2]:    id  x.radius_mean  x.texture_mean  x.perimeter_mean  x.area_mean  \
0    1         13.540         14.36         87.46         566.3
1    2         13.080         15.71         85.63         520.0
2    3          9.504         12.44         60.34         273.9
3    4         13.030         18.42         82.61         523.8
4    5          8.196         16.84         51.71         201.9

      x.smoothness_mean  x.compactness_mean  x.concavity_mean  \
0          0.09779          0.08129          0.06664
1          0.10750          0.12700          0.04568
2          0.10240          0.06492          0.02956
3          0.08983          0.03766          0.02562
4          0.08600          0.05943          0.01588

      x.concave_pts_mean  x.symmetry_mean  ...  x.texture_worst  \
0          0.047810          0.1885  ...          19.26
1          0.031100          0.1967  ...          20.49
2          0.020760          0.1815  ...          15.66
3          0.029230          0.1467  ...          22.81
4          0.005917          0.1769  ...          21.96

      x.perimeter_worst  x.area_worst  x.smoothness_worst  x.compactness_worst  \
0          99.70          711.2          0.14400          0.17730
```

1	96.09	630.5	0.13120	0.27760
2	65.13	314.9	0.13240	0.11480
3	84.46	545.9	0.09701	0.04619
4	57.26	242.2	0.12970	0.13570

	x.concavity_worst	x.concave_pts_worst	x.symmetry_worst \
0	0.23900	0.12880	0.2977
1	0.18900	0.07283	0.3184
2	0.08867	0.06227	0.2450
3	0.04833	0.05013	0.1987
4	0.06880	0.02564	0.3105

	x.fractal_dim_worst	diagnosis
0	0.07259	B
1	0.08183	B
2	0.07773	B
3	0.06169	B
4	0.07409	B

[5 rows x 32 columns]

```
[3]: # Drop Unnecessary Columns if present
if 'Unnamed: 32' in data.columns:
    data.drop('Unnamed: 32', axis=1, inplace=True)
if 'id' in data.columns:
    data.drop('id', axis=1, inplace=True)
```

```
[4]: # Encode Labels
data['diagnosis'] = data['diagnosis'].map({'M': 1, 'B': 0})
data['diagnosis']
```

```
[4]: 0      0
      1      0
      2      0
      3      0
      4      0
      ..
564    1
565    1
566    1
567    1
568    1
Name: diagnosis, Length: 569, dtype: int64
```

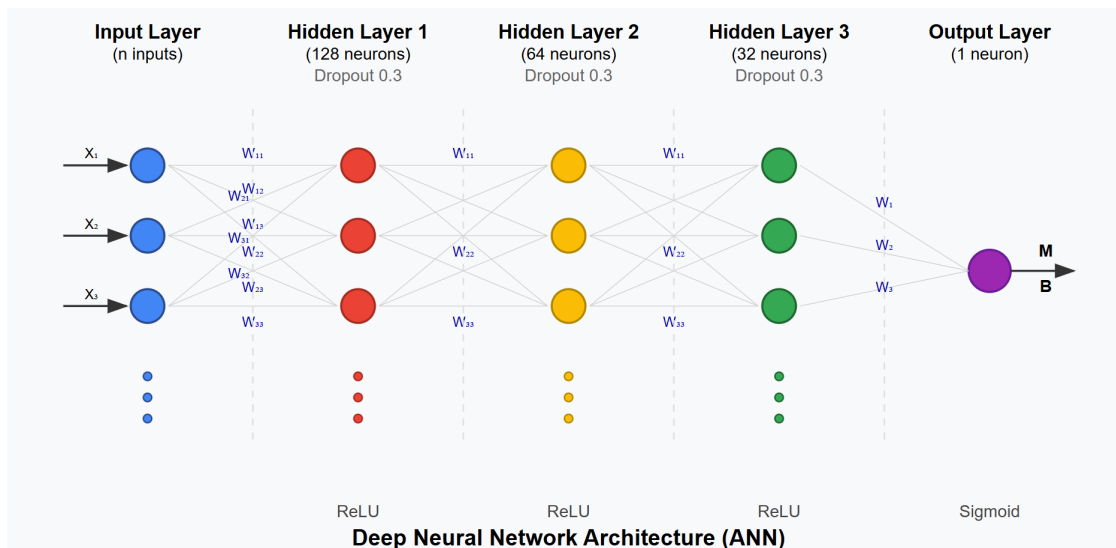
```
[5]: # Split features and labels
X = data.drop('diagnosis', axis=1)
y = data['diagnosis']
```

```
# Standardize the Data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
↳random_state=42)
```

```
[6]: # Build Deep Learning Model
model = Sequential([
    Dense(128, activation='relu', input_shape=(X_train.
↳shape[1],), name='Input_layer'),
    Dropout(0.3),
    Dense(64, activation='relu', name='First_hidden_layer'),
    Dropout(0.3),
    Dense(32, activation='relu', name='Second_hidden_layer'),
    Dropout(0.3),
    Dense(1, activation='sigmoid', name='Output_layer')
])
```

c:\Users\Uditya\Desktop\Deep Learning\Deep Learning for Beginner\venv\lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)



```
[7]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
Input_layer (Dense)	(None, 128)	3,968
dropout (Dropout)	(None, 128)	0
First_hidden_layer (Dense)	(None, 64)	8,256
dropout_1 (Dropout)	(None, 64)	0
Second_hidden_layer (Dense)	(None, 32)	2,080
dropout_2 (Dropout)	(None, 32)	0
Output_layer (Dense)	(None, 1)	33

Total params: 14,337 (56.00 KB)

Trainable params: 14,337 (56.00 KB)

Non-trainable params: 0 (0.00 B)

```
[8]: # Compile Model
model.compile(optimizer='adam', loss='binary_crossentropy',
             metrics=['accuracy'])

# Train Model
history = model.fit(X_train, y_train, epochs=100, batch_size=32,
                  validation_split=0.2, verbose=1)

# Evaluate Model
y_pred = (model.predict(X_test) > 0.5).astype("int32")
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))

print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Epoch 1/100

12/12 9s 88ms/step -

accuracy: 0.5576 - loss: 0.6892 - val_accuracy: 0.8901 - val_loss: 0.4203

Epoch 2/100

12/12 1s 29ms/step -

accuracy: 0.8557 - loss: 0.4311 - val_accuracy: 0.9231 - val_loss: 0.2509

Epoch 3/100

12/12 0s 26ms/step -
 accuracy: 0.9151 - loss: 0.2726 - val_accuracy: 0.9560 - val_loss: 0.1691
 Epoch 4/100
 12/12 0s 27ms/step -
 accuracy: 0.9625 - loss: 0.1675 - val_accuracy: 0.9560 - val_loss: 0.1234
 Epoch 5/100
 12/12 1s 23ms/step -
 accuracy: 0.9678 - loss: 0.1247 - val_accuracy: 0.9560 - val_loss: 0.0929
 Epoch 6/100
 12/12 0s 12ms/step -
 accuracy: 0.9721 - loss: 0.1096 - val_accuracy: 0.9670 - val_loss: 0.0738
 Epoch 7/100
 12/12 0s 18ms/step -
 accuracy: 0.9762 - loss: 0.1053 - val_accuracy: 0.9670 - val_loss: 0.0618
 Epoch 8/100
 12/12 0s 15ms/step -
 accuracy: 0.9643 - loss: 0.0888 - val_accuracy: 0.9670 - val_loss: 0.0610
 Epoch 9/100
 12/12 0s 20ms/step -
 accuracy: 0.9731 - loss: 0.1049 - val_accuracy: 0.9780 - val_loss: 0.0608
 Epoch 10/100
 12/12 1s 49ms/step -
 accuracy: 0.9823 - loss: 0.0766 - val_accuracy: 0.9780 - val_loss: 0.0735
 Epoch 11/100
 12/12 1s 43ms/step -
 accuracy: 0.9946 - loss: 0.0635 - val_accuracy: 0.9780 - val_loss: 0.0711
 Epoch 12/100
 12/12 0s 35ms/step -
 accuracy: 0.9761 - loss: 0.1296 - val_accuracy: 0.9780 - val_loss: 0.0694
 Epoch 13/100
 12/12 0s 18ms/step -
 accuracy: 0.9773 - loss: 0.0797 - val_accuracy: 0.9780 - val_loss: 0.0752
 Epoch 14/100
 12/12 0s 21ms/step -
 accuracy: 0.9872 - loss: 0.0378 - val_accuracy: 0.9780 - val_loss: 0.0846
 Epoch 15/100
 12/12 0s 18ms/step -
 accuracy: 0.9793 - loss: 0.0700 - val_accuracy: 0.9780 - val_loss: 0.0830
 Epoch 16/100
 12/12 0s 31ms/step -
 accuracy: 0.9862 - loss: 0.0586 - val_accuracy: 0.9780 - val_loss: 0.0817
 Epoch 17/100
 12/12 1s 47ms/step -
 accuracy: 0.9885 - loss: 0.0555 - val_accuracy: 0.9780 - val_loss: 0.0846
 Epoch 18/100
 12/12 0s 32ms/step -
 accuracy: 0.9892 - loss: 0.0325 - val_accuracy: 0.9780 - val_loss: 0.0890
 Epoch 19/100

12/12 0s 27ms/step -
 accuracy: 0.9914 - loss: 0.0512 - val_accuracy: 0.9780 - val_loss: 0.0916
 Epoch 20/100
 12/12 1s 42ms/step -
 accuracy: 0.9868 - loss: 0.0416 - val_accuracy: 0.9780 - val_loss: 0.0962
 Epoch 21/100
 12/12 1s 27ms/step -
 accuracy: 0.9898 - loss: 0.0376 - val_accuracy: 0.9780 - val_loss: 0.0941
 Epoch 22/100
 12/12 0s 24ms/step -
 accuracy: 0.9765 - loss: 0.0693 - val_accuracy: 0.9780 - val_loss: 0.0945
 Epoch 23/100
 12/12 0s 22ms/step -
 accuracy: 0.9850 - loss: 0.0402 - val_accuracy: 0.9780 - val_loss: 0.0966
 Epoch 24/100
 12/12 0s 22ms/step -
 accuracy: 0.9894 - loss: 0.0397 - val_accuracy: 0.9780 - val_loss: 0.1001
 Epoch 25/100
 12/12 0s 22ms/step -
 accuracy: 0.9902 - loss: 0.0357 - val_accuracy: 0.9670 - val_loss: 0.1054
 Epoch 26/100
 12/12 0s 24ms/step -
 accuracy: 0.9885 - loss: 0.0293 - val_accuracy: 0.9780 - val_loss: 0.1076
 Epoch 27/100
 12/12 0s 21ms/step -
 accuracy: 0.9864 - loss: 0.0553 - val_accuracy: 0.9780 - val_loss: 0.1079
 Epoch 28/100
 12/12 0s 20ms/step -
 accuracy: 0.9840 - loss: 0.0378 - val_accuracy: 0.9670 - val_loss: 0.1199
 Epoch 29/100
 12/12 1s 34ms/step -
 accuracy: 0.9824 - loss: 0.0322 - val_accuracy: 0.9670 - val_loss: 0.1356
 Epoch 30/100
 12/12 0s 19ms/step -
 accuracy: 0.9996 - loss: 0.0140 - val_accuracy: 0.9670 - val_loss: 0.1447
 Epoch 31/100
 12/12 0s 17ms/step -
 accuracy: 0.9914 - loss: 0.0231 - val_accuracy: 0.9670 - val_loss: 0.1355
 Epoch 32/100
 12/12 0s 17ms/step -
 accuracy: 0.9920 - loss: 0.0266 - val_accuracy: 0.9780 - val_loss: 0.1111
 Epoch 33/100
 12/12 0s 16ms/step -
 accuracy: 0.9962 - loss: 0.0153 - val_accuracy: 0.9670 - val_loss: 0.1270
 Epoch 34/100
 12/12 0s 17ms/step -
 accuracy: 0.9964 - loss: 0.0146 - val_accuracy: 0.9670 - val_loss: 0.1328
 Epoch 35/100

12/12 0s 17ms/step -
 accuracy: 0.9967 - loss: 0.0143 - val_accuracy: 0.9780 - val_loss: 0.1313
 Epoch 36/100
 12/12 0s 17ms/step -
 accuracy: 0.9887 - loss: 0.0199 - val_accuracy: 0.9780 - val_loss: 0.1283
 Epoch 37/100
 12/12 0s 17ms/step -
 accuracy: 0.9938 - loss: 0.0147 - val_accuracy: 0.9780 - val_loss: 0.1373
 Epoch 38/100
 12/12 0s 17ms/step -
 accuracy: 0.9994 - loss: 0.0073 - val_accuracy: 0.9780 - val_loss: 0.1410
 Epoch 39/100
 12/12 0s 28ms/step -
 accuracy: 0.9982 - loss: 0.0152 - val_accuracy: 0.9780 - val_loss: 0.1403
 Epoch 40/100
 12/12 1s 18ms/step -
 accuracy: 0.9937 - loss: 0.0167 - val_accuracy: 0.9780 - val_loss: 0.1188
 Epoch 41/100
 12/12 0s 20ms/step -
 accuracy: 0.9950 - loss: 0.0162 - val_accuracy: 0.9780 - val_loss: 0.1186
 Epoch 42/100
 12/12 0s 16ms/step -
 accuracy: 0.9978 - loss: 0.0118 - val_accuracy: 0.9780 - val_loss: 0.1261
 Epoch 43/100
 12/12 0s 21ms/step -
 accuracy: 0.9956 - loss: 0.0129 - val_accuracy: 0.9780 - val_loss: 0.1336
 Epoch 44/100
 12/12 0s 16ms/step -
 accuracy: 0.9974 - loss: 0.0077 - val_accuracy: 0.9780 - val_loss: 0.1365
 Epoch 45/100
 12/12 0s 23ms/step -
 accuracy: 0.9967 - loss: 0.0171 - val_accuracy: 0.9670 - val_loss: 0.1485
 Epoch 46/100
 12/12 0s 20ms/step -
 accuracy: 0.9959 - loss: 0.0111 - val_accuracy: 0.9780 - val_loss: 0.1488
 Epoch 47/100
 12/12 0s 16ms/step -
 accuracy: 0.9973 - loss: 0.0059 - val_accuracy: 0.9780 - val_loss: 0.1579
 Epoch 48/100
 12/12 0s 24ms/step -
 accuracy: 0.9872 - loss: 0.0250 - val_accuracy: 0.9780 - val_loss: 0.1496
 Epoch 49/100
 12/12 0s 15ms/step -
 accuracy: 0.9947 - loss: 0.0129 - val_accuracy: 0.9670 - val_loss: 0.1563
 Epoch 50/100
 12/12 0s 16ms/step -
 accuracy: 0.9967 - loss: 0.0102 - val_accuracy: 0.9780 - val_loss: 0.1654
 Epoch 51/100

12/12 0s 16ms/step -
 accuracy: 1.0000 - loss: 0.0054 - val_accuracy: 0.9780 - val_loss: 0.1721
 Epoch 52/100
 12/12 0s 17ms/step -
 accuracy: 0.9973 - loss: 0.0041 - val_accuracy: 0.9780 - val_loss: 0.1766
 Epoch 53/100
 12/12 0s 20ms/step -
 accuracy: 0.9986 - loss: 0.0057 - val_accuracy: 0.9670 - val_loss: 0.1790
 Epoch 54/100
 12/12 0s 18ms/step -
 accuracy: 0.9996 - loss: 0.0038 - val_accuracy: 0.9670 - val_loss: 0.1808
 Epoch 55/100
 12/12 0s 17ms/step -
 accuracy: 1.0000 - loss: 0.0060 - val_accuracy: 0.9670 - val_loss: 0.1762
 Epoch 56/100
 12/12 0s 18ms/step -
 accuracy: 1.0000 - loss: 0.0056 - val_accuracy: 0.9670 - val_loss: 0.1903
 Epoch 57/100
 12/12 0s 18ms/step -
 accuracy: 0.9996 - loss: 0.0030 - val_accuracy: 0.9560 - val_loss: 0.2012
 Epoch 58/100
 12/12 0s 13ms/step -
 accuracy: 0.9967 - loss: 0.0069 - val_accuracy: 0.9560 - val_loss: 0.2002
 Epoch 59/100
 12/12 0s 16ms/step -
 accuracy: 0.9991 - loss: 0.0049 - val_accuracy: 0.9670 - val_loss: 0.1851
 Epoch 60/100
 12/12 0s 13ms/step -
 accuracy: 0.9991 - loss: 0.0101 - val_accuracy: 0.9670 - val_loss: 0.1944
 Epoch 61/100
 12/12 0s 13ms/step -
 accuracy: 1.0000 - loss: 0.0049 - val_accuracy: 0.9560 - val_loss: 0.2376
 Epoch 62/100
 12/12 0s 15ms/step -
 accuracy: 0.9967 - loss: 0.0096 - val_accuracy: 0.9670 - val_loss: 0.2788
 Epoch 63/100
 12/12 0s 16ms/step -
 accuracy: 0.9974 - loss: 0.0040 - val_accuracy: 0.9670 - val_loss: 0.2756
 Epoch 64/100
 12/12 0s 15ms/step -
 accuracy: 1.0000 - loss: 0.0036 - val_accuracy: 0.9560 - val_loss: 0.2298
 Epoch 65/100
 12/12 0s 16ms/step -
 accuracy: 0.9978 - loss: 0.0063 - val_accuracy: 0.9670 - val_loss: 0.2481
 Epoch 66/100
 12/12 0s 18ms/step -
 accuracy: 0.9969 - loss: 0.0052 - val_accuracy: 0.9670 - val_loss: 0.2811
 Epoch 67/100

12/12 0s 13ms/step -
accuracy: 1.0000 - loss: 0.0032 - val_accuracy: 0.9560 - val_loss: 0.2951
Epoch 68/100

12/12 0s 14ms/step -
accuracy: 1.0000 - loss: 0.0018 - val_accuracy: 0.9560 - val_loss: 0.2951
Epoch 69/100

12/12 0s 23ms/step -
accuracy: 0.9986 - loss: 0.0109 - val_accuracy: 0.9670 - val_loss: 0.2263
Epoch 70/100

12/12 0s 13ms/step -
accuracy: 1.0000 - loss: 0.0015 - val_accuracy: 0.9670 - val_loss: 0.2079
Epoch 71/100

12/12 0s 15ms/step -
accuracy: 0.9919 - loss: 0.0143 - val_accuracy: 0.9670 - val_loss: 0.3645
Epoch 72/100

12/12 0s 16ms/step -
accuracy: 0.9982 - loss: 0.0067 - val_accuracy: 0.9780 - val_loss: 0.1734
Epoch 73/100

12/12 0s 16ms/step -
accuracy: 0.9973 - loss: 0.0153 - val_accuracy: 0.9670 - val_loss: 0.1980
Epoch 74/100

12/12 0s 15ms/step -
accuracy: 1.0000 - loss: 0.0035 - val_accuracy: 0.9670 - val_loss: 0.2791
Epoch 75/100

12/12 0s 27ms/step -
accuracy: 1.0000 - loss: 6.9483e-04 - val_accuracy: 0.9670 - val_loss: 0.3047
Epoch 76/100

12/12 0s 15ms/step -
accuracy: 1.0000 - loss: 0.0032 - val_accuracy: 0.9670 - val_loss: 0.3255
Epoch 77/100

12/12 0s 17ms/step -
accuracy: 0.9989 - loss: 0.0096 - val_accuracy: 0.9670 - val_loss: 0.2980
Epoch 78/100

12/12 0s 13ms/step -
accuracy: 0.9989 - loss: 0.0071 - val_accuracy: 0.9670 - val_loss: 0.3283
Epoch 79/100

12/12 0s 15ms/step -
accuracy: 0.9947 - loss: 0.0065 - val_accuracy: 0.9670 - val_loss: 0.3487
Epoch 80/100

12/12 0s 23ms/step -
accuracy: 0.9931 - loss: 0.0333 - val_accuracy: 0.9780 - val_loss: 0.1954
Epoch 81/100

12/12 0s 15ms/step -
accuracy: 0.9890 - loss: 0.0190 - val_accuracy: 0.9780 - val_loss: 0.1776
Epoch 82/100

12/12 0s 20ms/step -
accuracy: 1.0000 - loss: 0.0029 - val_accuracy: 0.9780 - val_loss: 0.1733
Epoch 83/100

12/12 0s 19ms/step -
 accuracy: 0.9890 - loss: 0.0134 - val_accuracy: 0.9780 - val_loss: 0.1599
 Epoch 84/100
 12/12 0s 13ms/step -
 accuracy: 1.0000 - loss: 0.0017 - val_accuracy: 0.9780 - val_loss: 0.1628
 Epoch 85/100
 12/12 0s 16ms/step -
 accuracy: 0.9986 - loss: 0.0038 - val_accuracy: 0.9670 - val_loss: 0.1821
 Epoch 86/100
 12/12 0s 22ms/step -
 accuracy: 1.0000 - loss: 0.0016 - val_accuracy: 0.9670 - val_loss: 0.2046
 Epoch 87/100
 12/12 0s 15ms/step -
 accuracy: 0.9982 - loss: 0.0054 - val_accuracy: 0.9670 - val_loss: 0.2157
 Epoch 88/100
 12/12 0s 21ms/step -
 accuracy: 1.0000 - loss: 0.0017 - val_accuracy: 0.9670 - val_loss: 0.2337
 Epoch 89/100
 12/12 0s 16ms/step -
 accuracy: 1.0000 - loss: 0.0030 - val_accuracy: 0.9670 - val_loss: 0.2527
 Epoch 90/100
 12/12 0s 13ms/step -
 accuracy: 1.0000 - loss: 0.0020 - val_accuracy: 0.9670 - val_loss: 0.2590
 Epoch 91/100
 12/12 0s 16ms/step -
 accuracy: 1.0000 - loss: 0.0017 - val_accuracy: 0.9670 - val_loss: 0.2681
 Epoch 92/100
 12/12 0s 18ms/step -
 accuracy: 1.0000 - loss: 0.0017 - val_accuracy: 0.9670 - val_loss: 0.2853
 Epoch 93/100
 12/12 0s 19ms/step -
 accuracy: 1.0000 - loss: 0.0040 - val_accuracy: 0.9670 - val_loss: 0.2988
 Epoch 94/100
 12/12 0s 17ms/step -
 accuracy: 1.0000 - loss: 0.0031 - val_accuracy: 0.9670 - val_loss: 0.3107
 Epoch 95/100
 12/12 0s 17ms/step -
 accuracy: 0.9994 - loss: 0.0013 - val_accuracy: 0.9670 - val_loss: 0.3193
 Epoch 96/100
 12/12 0s 16ms/step -
 accuracy: 1.0000 - loss: 0.0016 - val_accuracy: 0.9670 - val_loss: 0.3429
 Epoch 97/100
 12/12 0s 14ms/step -
 accuracy: 1.0000 - loss: 0.0025 - val_accuracy: 0.9670 - val_loss: 0.3216
 Epoch 98/100
 12/12 0s 22ms/step -
 accuracy: 1.0000 - loss: 0.0015 - val_accuracy: 0.9780 - val_loss: 0.2150
 Epoch 99/100

12/12 0s 13ms/step -
accuracy: 0.9958 - loss: 0.0209 - val_accuracy: 0.9670 - val_loss: 0.2285
Epoch 100/100

12/12 0s 14ms/step -
accuracy: 0.9982 - loss: 0.0088 - val_accuracy: 0.9780 - val_loss: 0.2131
4/4 0s 38ms/step

Confusion Matrix:

```
[[70  1]
 [ 3 40]]
```

Classification Report:


	precision	recall	f1-score	support
0	0.96	0.99	0.97	71
1	0.98	0.93	0.95	43
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

```
[ ]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
# Predict on test data
y_pred_probs = model.predict(X_test)
y_pred = (y_pred_probs > 0.5).astype("int32")

# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)

disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Benign",
↪ "Malignant"])
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix")
plt.savefig("results/confusion_metrics.png")
plt.show()
```

4/4 0s 27ms/step



```
[10]: import os

# Create results folder if not exists
os.makedirs("results", exist_ok=True)


# Plot Accuracy and Loss
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Val Accuracy')
plt.title('Model Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('Model Loss')
plt.legend()
```

```
plt.tight_layout()

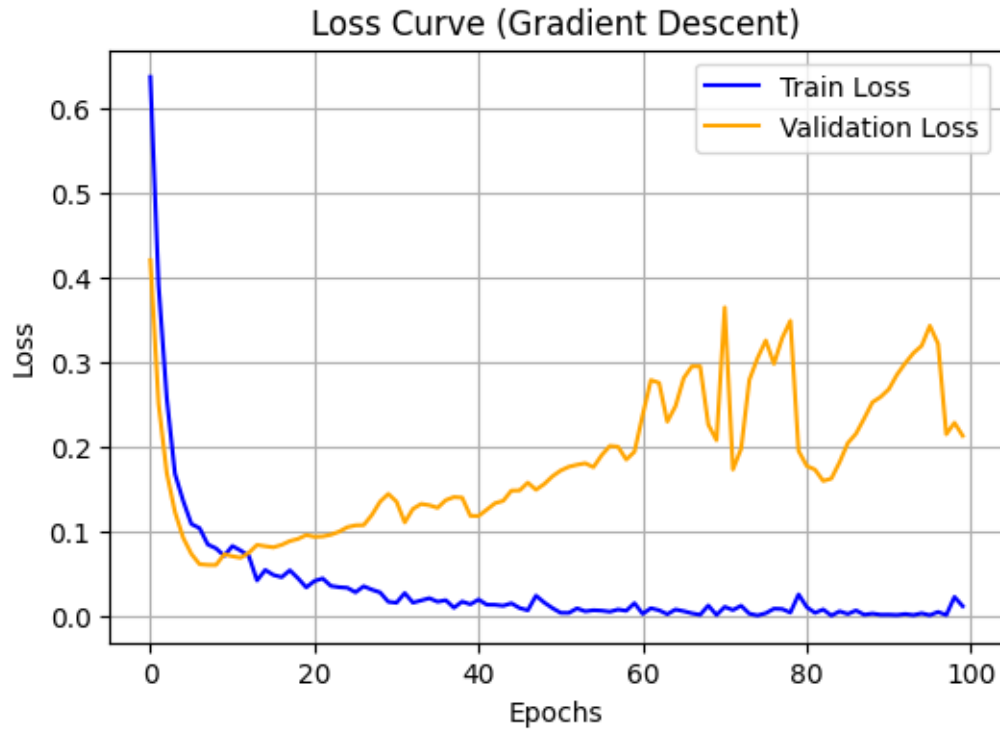
# Save the plot
plt.savefig("results/training_metrics.png")
plt.show()
```



```
[11]: # Plotting only Loss (Gradient Descent Visualization)
import os
os.makedirs("results", exist_ok=True)

plt.figure(figsize=(6, 4))
plt.plot(history.history['loss'], label='Train Loss', color='blue')
plt.plot(history.history['val_loss'], label='Validation Loss', color='orange')
plt.title('Loss Curve (Gradient Descent)')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)

plt.savefig("results/loss_gradient_descent.png")
plt.show()
```



```
[12]: # Prepare Dataset
train_dataset = tf.data.Dataset.from_tensor_slices((X_train, y_train)).batch(16)

# Optimizer & Loss
loss_fn = tf.keras.losses.BinaryCrossentropy()
optimizer = tf.keras.optimizers.Adam()

# For Tracking
gradient_norms = {layer.name: [] for layer in model.layers if len(layer.
    trainable_variables) > 0}
losses = []

# Training Loop
epochs = 100
for epoch in range(epochs):
    print(f"\nEpoch {epoch+1}/{epochs}")
    epoch_losses = []

    for step, (x_batch, y_batch) in enumerate(train_dataset):
        with tf.GradientTape() as tape:
            logits = model(x_batch, training=True)
            loss_value = loss_fn(y_batch, logits)
```

```

grads = tape.gradient(loss_value, model.trainable_variables)
optimizer.apply_gradients(zip(grads, model.trainable_variables))
epoch_losses.append(loss_value.numpy())

# Save gradient norms
idx = 0
for layer in model.layers:
    if len(layer.trainable_variables) > 0:
        for var in layer.trainable_variables:
            grad = grads[idx]
            norm = tf.norm(grad).numpy() if grad is not None else 0
            gradient_norms[layer.name].append(norm)
            idx += 1

# Average loss of epoch
epoch_loss = np.mean(epoch_losses)
losses.append(epoch_loss)
print(f"Loss: {epoch_loss:.4f}")

# Plotting
plt.figure(figsize=(14, 5), dpi=200)

# Loss vs Epoch
plt.subplot(1, 2, 1)
plt.plot(range(1, epochs+1), losses, marker='o', color='blue')
plt.title("Loss vs Epochs")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.grid(True)

# Gradient Norm vs Epoch (per layer)
plt.subplot(1, 2, 2)
steps_per_epoch = len(train_dataset)

for layer_name, norms in gradient_norms.items():
    norms = np.array(norms)
    avg_per_epoch = [np.mean(norms[i*steps_per_epoch:(i+1)*steps_per_epoch])]
    for i in range(epochs)]
    plt.plot(range(1, epochs+1), avg_per_epoch, label=layer_name)

plt.title("Gradient Norm vs Epochs (per layer)")
plt.xlabel("Epochs")
plt.ylabel("Avg Gradient Norm")
plt.legend()
plt.grid(True)

plt.tight_layout()

```

```
plt.savefig("results/gradient_norm_per_epoch.png")  
plt.show()
```

Epoch 1/100
Loss: 0.0767

Epoch 2/100
Loss: 0.0339

Epoch 3/100
Loss: 0.1016

Epoch 4/100
Loss: 0.0270

Epoch 5/100
Loss: 0.0352

Epoch 6/100
Loss: 0.0327

Epoch 7/100
Loss: 0.0108

Epoch 8/100
Loss: 0.0140

Epoch 9/100
Loss: 0.0108

Epoch 10/100
Loss: 0.0175

Epoch 11/100
Loss: 0.0129

Epoch 12/100
Loss: 0.0269

Epoch 13/100
Loss: 0.0159

Epoch 14/100
Loss: 0.0097

Epoch 15/100
Loss: 0.0078

Epoch 16/100
Loss: 0.0082

Epoch 17/100
Loss: 0.0052

Epoch 18/100
Loss: 0.0094

Epoch 19/100
Loss: 0.0046

Epoch 20/100
Loss: 0.0038

Epoch 21/100
Loss: 0.0072

Epoch 22/100
Loss: 0.0057

Epoch 23/100
Loss: 0.0047

Epoch 24/100
Loss: 0.0037

Epoch 25/100
Loss: 0.0011

Epoch 26/100
Loss: 0.0092

Epoch 27/100
Loss: 0.0328

Epoch 28/100
Loss: 0.0046

Epoch 29/100
Loss: 0.0137

Epoch 30/100
Loss: 0.0038

Epoch 31/100
Loss: 0.0062

Epoch 32/100
Loss: 0.0043

Epoch 33/100
Loss: 0.0052

Epoch 34/100
Loss: 0.0071

Epoch 35/100
Loss: 0.0015

Epoch 36/100
Loss: 0.0031

Epoch 37/100
Loss: 0.0018

Epoch 38/100
Loss: 0.0015

Epoch 39/100
Loss: 0.0021

Epoch 40/100
Loss: 0.0044

Epoch 41/100
Loss: 0.0020

Epoch 42/100
Loss: 0.0041

Epoch 43/100
Loss: 0.0023

Epoch 44/100
Loss: 0.0025

Epoch 45/100
Loss: 0.0008

Epoch 46/100
Loss: 0.0013

Epoch 47/100
Loss: 0.0051

Epoch 48/100
Loss: 0.0006

Epoch 49/100
Loss: 0.0010

Epoch 50/100
Loss: 0.0037

Epoch 51/100
Loss: 0.0039

Epoch 52/100
Loss: 0.0049

Epoch 53/100
Loss: 0.0011

Epoch 54/100
Loss: 0.0048

Epoch 55/100
Loss: 0.0073

Epoch 56/100
Loss: 0.0010

Epoch 57/100
Loss: 0.0098

Epoch 58/100
Loss: 0.0044

Epoch 59/100
Loss: 0.0005

Epoch 60/100
Loss: 0.0025

Epoch 61/100
Loss: 0.0004

Epoch 62/100
Loss: 0.0178

Epoch 63/100
Loss: 0.0027

Epoch 64/100
Loss: 0.0067

Epoch 65/100
Loss: 0.0064

Epoch 66/100
Loss: 0.0019

Epoch 67/100
Loss: 0.0055

Epoch 68/100
Loss: 0.0085

Epoch 69/100
Loss: 0.0036

Epoch 70/100
Loss: 0.0095

Epoch 71/100
Loss: 0.0041

Epoch 72/100
Loss: 0.0115

Epoch 73/100
Loss: 0.0039

Epoch 74/100
Loss: 0.0004

Epoch 75/100
Loss: 0.0011

Epoch 76/100
Loss: 0.0011

Epoch 77/100
Loss: 0.0012

Epoch 78/100
Loss: 0.0007

Epoch 79/100
Loss: 0.0030

Epoch 80/100
Loss: 0.0118

Epoch 81/100
Loss: 0.0015

Epoch 82/100
Loss: 0.0041

Epoch 83/100
Loss: 0.0012

Epoch 84/100
Loss: 0.0008

Epoch 85/100
Loss: 0.0013

Epoch 86/100
Loss: 0.0004

Epoch 87/100
Loss: 0.0012

Epoch 88/100
Loss: 0.0004

Epoch 89/100
Loss: 0.0003

Epoch 90/100
Loss: 0.0004

Epoch 91/100
Loss: 0.0004

Epoch 92/100
Loss: 0.0008

Epoch 93/100
Loss: 0.0007

Epoch 94/100
Loss: 0.0002

Epoch 95/100
Loss: 0.0004

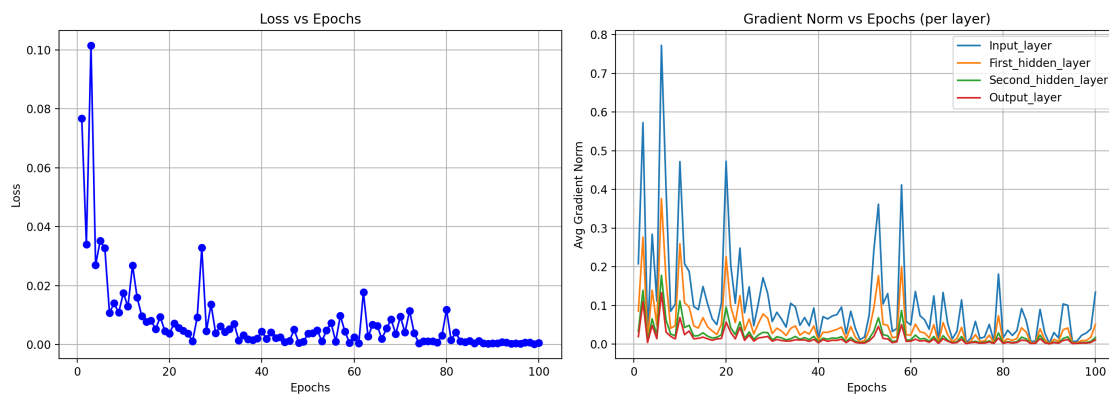
Epoch 96/100
Loss: 0.0003

Epoch 97/100
Loss: 0.0007

Epoch 98/100
Loss: 0.0007

Epoch 99/100
Loss: 0.0002


Epoch 100/100
Loss: 0.0006



```
[13]: # Create result directory if not exists
import os
os.makedirs("results", exist_ok=True)

# Plot gradient norms for each layer
plt.figure(figsize=(10, 6), dpi= 200)
for layer_name, norms in gradient_norms.items():
    plt.plot(norms, label=layer_name)

plt.title("Gradient Norms per Layer across Batches")
plt.xlabel("Batch Steps")
plt.ylabel("Gradient Norm")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig("results/gradient_flow_per_layer.png")
plt.show()
```



0.1 Binary Cross Entropy Loss

core of **gradient descent** update:

$$W_{t+1} = W_t - \eta \cdot \frac{\partial L}{\partial W}$$

core of **gradient descent** update:

$$W_{t+1} = W_t - \eta \cdot \frac{\partial L}{\partial W}$$

core of gradient descent update:

$$W_{t+1} = W_t - \eta \cdot \frac{\partial L}{\partial W}$$

```
[16]: loss_fn = tf.keras.losses.BinaryCrossentropy()
optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)

# To store gradient norms
gradient_norms = []
loss_values = []

# Convert data to tensors
X_tensor = tf.convert_to_tensor(X_train, dtype=tf.float32)
y_tensor = tf.convert_to_tensor(y_train.values.reshape(-1, 1), dtype=tf.float32)

# Training Loop for 100 epochs
for epoch in range(100):
    with tf.GradientTape() as tape:
        predictions = model(X_tensor, training=True)
        loss_value = loss_fn(y_tensor, predictions)

    # Compute gradients
    grads = tape.gradient(loss_value, model.trainable_weights)

    # Record gradient norms
    loss_values.append(loss_value.numpy())
    grad_norm = np.mean([tf.norm(g).numpy() for g in grads if g is not None])
    gradient_norms.append(grad_norm)

    # Apply gradients
    optimizer.apply_gradients(zip(grads, model.trainable_weights))

    print(f"Epoch {epoch+1}: Loss = {loss_value:.4f}, Avg. Grad Norm = {grad_norm:.4f}")

# Plot gradient norms
plt.figure(figsize=(12, 7), dpi=250)
plt.plot(gradient_norms, marker='o', color='green')
plt.title("Average Gradient Norm per Epoch")
plt.xlabel("Epoch")
plt.ylabel("Gradient Norm")
plt.grid(True)
```




```
plt.tight_layout()
plt.savefig("results/gradient_flow.png")
plt.show()
```

```
Epoch 1: Loss = 0.0001, Avg. Grad Norm = 0.0004
Epoch 2: Loss = 0.0002, Avg. Grad Norm = 0.0011
Epoch 3: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 4: Loss = 0.0001, Avg. Grad Norm = 0.0004
Epoch 5: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 6: Loss = 0.0001, Avg. Grad Norm = 0.0002
Epoch 7: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 8: Loss = 0.0000, Avg. Grad Norm = 0.0003
Epoch 9: Loss = 0.0005, Avg. Grad Norm = 0.0096
Epoch 10: Loss = 0.0001, Avg. Grad Norm = 0.0004
Epoch 11: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 12: Loss = 0.0001, Avg. Grad Norm = 0.0020
Epoch 13: Loss = 0.0003, Avg. Grad Norm = 0.0047
Epoch 14: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 15: Loss = 0.0000, Avg. Grad Norm = 0.0004
Epoch 16: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 17: Loss = 0.0000, Avg. Grad Norm = 0.0003
Epoch 18: Loss = 0.0001, Avg. Grad Norm = 0.0009
Epoch 19: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 20: Loss = 0.0008, Avg. Grad Norm = 0.0170
Epoch 21: Loss = 0.0001, Avg. Grad Norm = 0.0008
Epoch 22: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 23: Loss = 0.0006, Avg. Grad Norm = 0.0097
Epoch 24: Loss = 0.0000, Avg. Grad Norm = 0.0003
Epoch 25: Loss = 0.0000, Avg. Grad Norm = 0.0007
Epoch 26: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 27: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 28: Loss = 0.0003, Avg. Grad Norm = 0.0045
Epoch 29: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 30: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 31: Loss = 0.0002, Avg. Grad Norm = 0.0021
Epoch 32: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 33: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 34: Loss = 0.0005, Avg. Grad Norm = 0.0068
Epoch 35: Loss = 0.0001, Avg. Grad Norm = 0.0010
Epoch 36: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 37: Loss = 0.0005, Avg. Grad Norm = 0.0063
Epoch 38: Loss = 0.0001, Avg. Grad Norm = 0.0026
Epoch 39: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 40: Loss = 0.0002, Avg. Grad Norm = 0.0018
Epoch 41: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 42: Loss = 0.0000, Avg. Grad Norm = 0.0003
Epoch 43: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 44: Loss = 0.0000, Avg. Grad Norm = 0.0002
```

Epoch 45: Loss = 0.0006, Avg. Grad Norm = 0.0076
Epoch 46: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 47: Loss = 0.0001, Avg. Grad Norm = 0.0005
Epoch 48: Loss = 0.0003, Avg. Grad Norm = 0.0093
Epoch 49: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 50: Loss = 0.0001, Avg. Grad Norm = 0.0012
Epoch 51: Loss = 0.0023, Avg. Grad Norm = 0.0245
Epoch 52: Loss = 0.0001, Avg. Grad Norm = 0.0002
Epoch 53: Loss = 0.0000, Avg. Grad Norm = 0.0003
Epoch 54: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 55: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 56: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 57: Loss = 0.0006, Avg. Grad Norm = 0.0097
Epoch 58: Loss = 0.0000, Avg. Grad Norm = 0.0003
Epoch 59: Loss = 0.0002, Avg. Grad Norm = 0.0011
Epoch 60: Loss = 0.0001, Avg. Grad Norm = 0.0011
Epoch 61: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 62: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 63: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 64: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 65: Loss = 0.0019, Avg. Grad Norm = 0.0196
Epoch 66: Loss = 0.0000, Avg. Grad Norm = 0.0003
Epoch 67: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 68: Loss = 0.0056, Avg. Grad Norm = 0.0620
Epoch 69: Loss = 0.0028, Avg. Grad Norm = 0.0360
Epoch 70: Loss = 0.0001, Avg. Grad Norm = 0.0003
Epoch 71: Loss = 0.0001, Avg. Grad Norm = 0.0008
Epoch 72: Loss = 0.0001, Avg. Grad Norm = 0.0004
Epoch 73: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 74: Loss = 0.0001, Avg. Grad Norm = 0.0009
Epoch 75: Loss = 0.0001, Avg. Grad Norm = 0.0003
Epoch 76: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 77: Loss = 0.0027, Avg. Grad Norm = 0.0253
Epoch 78: Loss = 0.0001, Avg. Grad Norm = 0.0010
Epoch 79: Loss = 0.0001, Avg. Grad Norm = 0.0008
Epoch 80: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 81: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 82: Loss = 0.0006, Avg. Grad Norm = 0.0048
Epoch 83: Loss = 0.0014, Avg. Grad Norm = 0.0178
Epoch 84: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 85: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 86: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 87: Loss = 0.0000, Avg. Grad Norm = 0.0002
Epoch 88: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 89: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 90: Loss = 0.0000, Avg. Grad Norm = 0.0001
Epoch 91: Loss = 0.0000, Avg. Grad Norm = 0.0000
Epoch 92: Loss = 0.0001, Avg. Grad Norm = 0.0006

Epoch 93: Loss = 0.0000, Avg. Grad Norm = 0.0002
 Epoch 94: Loss = 0.0001, Avg. Grad Norm = 0.0008
 Epoch 95: Loss = 0.0000, Avg. Grad Norm = 0.0001
 Epoch 96: Loss = 0.0000, Avg. Grad Norm = 0.0000
 Epoch 97: Loss = 0.0000, Avg. Grad Norm = 0.0002
 Epoch 98: Loss = 0.0000, Avg. Grad Norm = 0.0000
 Epoch 99: Loss = 0.0000, Avg. Grad Norm = 0.0003
 Epoch 100: Loss = 0.0000, Avg. Grad Norm = 0.0000



```
[17]: plt.figure(figsize=(12, 6), dpi=250)
plt.plot(loss_values, marker='o', color='orange')
plt.title("Loss vs Epoch")
plt.xlabel("Epoch")
plt.ylabel("Binary Crossentropy Loss")
plt.grid(True)
plt.tight_layout()
plt.savefig("results/loss_vs_epoch.png")
plt.show()
```

